


# Support Vector Machine


# Support Vector Machine - Classification (SVM)

- A Support Vector Machine (SVM) performs classification by finding the hyperplane that maximizes the margin between the two classes.
- The vectors (cases) that define the hyperplane are the support vectors.




# How does SVM work?

- Let's imagine we have two classes: red and blue, and our data has two features:  $x$  and  $y$ .
- We want a classifier that, given a pair of  $(x,y)$  coordinates, outputs if it's either red or blue. We plot our already labeled training data on a plane:




# How does SVM work?

- A SVM takes these data points and outputs the hyperplane that best separates the classes.
- This line is the decision boundary: anything that falls to one side of it we will classify as blue, and anything that falls to the other as red.




# How does SVM work?

- But, what exactly is the best hyperplane?
- For SVM, it's the one that maximizes the margins from both classes.
- In other words: the hyperplane (remember it's a line in this case) whose distance to the nearest element of each tag is the largest.
- Support vectors are the data points nearest to the hyperplane.




# How does SVM work?


- You can think of a hyperplane as a line that linearly separates and classifies a set of data.
- Intuitively, the further from the hyperplane our data points lie, the more confident we are that they have been correctly classified.
- We therefore want our data points to be as far away from the hyperplane as possible.



# How does SVM work?




# How does SVM work?






# How does SVM work?




# Some definitions

- Margin of Separation ( $d$ ): the separation between the hyperplane and the closest data point.
- Optimal Hyperplane (maximal margin): the particular hyperplane for which the margin of separation  $d$  is maximized.
- We want a classifier (linear separator) with as big a margin as possible.





# Good vs Bad Classifiers

- In this case, finding a line that passes between the red and green clusters, is a good algorithm.



# Good vs Bad Classifiers

- The line here is our separating boundary.
- The figure shows two possible classifiers for our problem.





# Good vs Bad Classifiers

- Here's an interesting question: both lines above separate the red and green clusters.
- Is there a good reason to choose one over another?
- The first line seems a bit “skewed.”
- Sure, it separates the training data perfectly, but if it sees a test point that's a little farther out from the clusters, there is a good chance it would get the label wrong.
- The second line doesn't have this problem. For example, look at the test points shown as squares and the labels assigned by the classifiers in the figure next.


# Good vs Bad Classifiers

- The second line stays as far away as possible from both the clusters while getting the training data separation right.
- By being right in the middle of the two clusters, it is less “risky,” gives the data distributions for each class some scope, and thus generalizes well on test data.



# Good vs Bad Classifiers


- SVMs try to find the second kind of line.
- Here's the second line shown with the support vectors:
  - points with black edges (there are two of them) and
  - the margin (the shaded region).





# Allowing for Errors

- We looked at the easy case of perfectly linearly separable data in the last section.
- Real-world data is, however, typically messy.
- You will almost always have a few instances that a linear classifier can't get right.
- Here's an example of such data









# Allowing for Errors

- You can provide a parameter called “C” to your SVM; this allows you to dictate the tradeoff between
  - Having a wide margin.
  - Correctly classifying training data. A higher value of C implies you want lesser errors on the training data.
- You get better classification of training data at the expense of a wide margin.
- Plots on next slide show how the classifier and the margin vary as we increase the value of C

# Allowing for Errors

- Note how the line “tilts” as we increase the value of  $C$ .
- Since this is a trade-off, note how the width of the margin shrinks as we increase the value of  $C$ .



# Pros & Cons of Support Vector Machines

- Pros
  - Accuracy
  - Works well on smaller cleaner datasets
- Cons
  - Isn't suited to larger datasets as the training time with SVMs can be high
  - Less effective on noisier datasets with overlapping classes

Thanks